

Reinforcement Learning Lab

Lesson 4: Temporal Difference Methods

Davide Corsi and Alberto Castellini

University of Verona
email: davide.corsi@univr.it

Academic Year 2022-23



UNIVERSITÀ
di VERONA
Dipartimento
di **INFORMATICA**

Environment Setup

The first step for the setup of the laboratory environment is to update the repository and load the **miniconda** environment.

- Update the repository of the lab:

```
cd RL-Lab  
git stash  
git pull  
git stash pop
```

- Activate the *miniconda* environment:

```
conda activate rl-lab
```

Safe Procedure

Always back up the previous lessons' solutions before executing the repository update.

Today Assignment

In today's lesson, we implement the **Q-Learning** and **SARSA** algorithms in Python. In particular, the file to complete is:

`RL-Lab/lessons/lesson_4_code.py`

Inside the file, two functions are partially implemented. The objective of this lesson is to complete them.

- **def QLearning()**
- **def SARSA()**

Expected results can be found in:

`RL-Lab/results/lesson_4_results.txt`

Q-Learning

Require: *environment* $[A, S]$, *problem*, *episodes*, α, γ , *expl_func*, *expl_param*

Ensure: *policy*, *rewards*, *lengths*

1: $\forall a \in A, \forall s \in S$ initialize $Q(s, a)$ arbitrarily

2: $\text{rewards}, \text{lengths} \leftarrow [0, \dots, 0]$

▷ Null vectors of length *episodes*

3: **for** $i \leftarrow 0$ **to** *episodes* **do**

4: Initialize s

5: **repeat**

6: $a \leftarrow \text{EXPL_FUNC}(Q, s, \text{expl_param})$

7: $s', r \leftarrow \text{take action } a \text{ from state } s$

▷ Act and observe

8: $Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma \max_{a' \in A_s} Q(s', a') - Q(s, a))$

▷ TD

9: $s \leftarrow s'$

10: **until** s is terminal

11: Update *rewards*, *lengths*

12: $\pi \leftarrow [0, \dots, 0]$

▷ Null vector of length $|S|$

13: **for each** s **in** S **do**

▷ Extract policy

14: $\pi_s \leftarrow \operatorname{argmax}_{a \in A_s} Q(s, a)$

15: **return** $\pi, \text{rewards}, \text{lengths}$

SARSA

Require: *environment* $[A, S]$, *problem*, *episodes*, α, γ , *expl_func*, *expl_param*

Ensure: *policy*, *rewards*, *lengths*

- 1: $\forall a \in A, \forall s \in S$ initialize $Q(s, a)$ arbitrarily
- 2: $\text{rewards}, \text{lengths} \leftarrow [0, \dots, 0]$ ▷ Null vectors of length *episodes*
- 3: **for** $i \leftarrow 0$ **to** *episodes* **do**
- 4: Initialize s
- 5: $a \leftarrow \text{EXPL_FUNC}(Q, s, \text{expl_param})$
- 6: **repeat**
- 7: $s', r \leftarrow$ take action a from state s ▷ Act and observe
- 8: $a' \leftarrow \text{EXPL_FUNC}(Q, s', \text{expl_param})$
- 9: $Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma Q(s', a') - Q(s, a))$ ▷ TD
- 10: $s \leftarrow s'$
- 11: $a \leftarrow a'$
- 12: **until** s is terminal
- 13: Update *rewards*, *lengths*
- 14: $\pi \leftarrow [0, \dots, 0]$ ▷ Null vector of length $|S|$
- 15: **for each** s **in** S **do** ▷ Extract policy
- 16: $\pi_s \leftarrow \underset{a \in A_s}{\operatorname{argmax}} Q(s, a)$