

Reinforcement Learning Lab

Lesson 8: Deep Q-Networks

Davide Corsi and Alberto Castellini

University of Verona
email: davide.corsi@univr.it

Academic Year 2022-23



UNIVERSITÀ
di VERONA
Dipartimento
di **INFORMATICA**

Environment Setup

The first step for the setup of the laboratory environment is to update the repository and load the **miniconda** environment.

- Update the repository of the lab:

```
cd RL-Lab
git stash
git pull
git stash pop
```


- Activate the *miniconda* environment:

```
conda activate rl-lab
```

- Install *Gymnasium* environments:

```
pip install gymnasium
```

Environment: CartPole



- A pole is attached by an unactuated joint to a cart, which moves along a frictionless track. The pendulum starts upright, and the goal is to prevent it from falling over. A **reward** of $+1$ is provided for every timestep that the pole remains upright.
- The **state** of the environment is represented as a tuple of 4 values: *Cart Position*, *Cart Velocity*, *Pole Angle*, and *Pole Angular Velocity*.
- The **actions** allowed in the environment are 2: *action 0 (push cart to left)* and *action 1 (push cart to right)*.

Today Assignment

In today's lesson, we will implement the **Deep Q-Network (DQN)** algorithm to solve the CartPole problem. In particular, the file to complete is:

`RL-Lab/lessons/lesson_8_code.py`

Inside the file, two python functions are partially implemented. The objective of this lesson is to complete it.

- **def training_loop()**
- **def DQNUpdate()**

Expected results can be found in:

`RL-Lab/results/lesson_8_results.txt`

Suggestions and Code Snippets

- 1 In today's lesson, the code is already partially implemented. Some functions related to Numpy and Matplotlib should not be modified. All the *entry points* for your code are marked with the keyword **TODO**.
- 2 Some methods of **Gymnasium** (the *new* Gym version) can be slightly different with respect to DangerousGridworld. Following are some snippets for the most important functions:

```
# Generation and Reset of the environment
env = gymnasium.make( "CartPole-v1" )
state = env.reset()[0]
# To generate a random action
action = env.action_space.sample()
# The updated 'step' function
next_state, reward, terminated, truncated, info = env.step(action)
done = terminated or truncated
```

training_loop()

Require: *environment, neural_network, trials, expl_param, score_queue*

Ensure: *neural_network, score_queue*

- 1: initialize the experience buffer
 - 2: initialize the score queue
 - 3: **for** $i \leftarrow 0$ **to** *epochs* **do**
 - 4: initialize *s* observe current state
 - 5: **repeat**
 - 6: Select and execute action *a*
 - 7: Observe new state s' and receive immediate reward *r*
 - 8: Add (*s*, *a*, s' , *r*) to experience buffer
 - 9: DQNUPLICATE(*neural_network*, *buffer*)
 - 10: update state $s \leftarrow s'$
 - 11: **until** *s* is terminal
 - 12: update *score_queue*
 - 13:
 - 14: **return** *score_queue*
- ▷ A fixed size queue
 - ▷ An infinite size queue
 - ▷ ϵ -greedy approach
 - ▷ Call the training function

DQNUpdate()

This function updates the neural network weights according to the formula:

$$w_{t+1} = w_t - learning_rate * gradient(R_{t+1} + \gamma \max_a (\hat{Q}_{pi}(S_{t+1}, a, w_t)) - \hat{Q}_{pi}(S_t, A_t, w_t))$$

Require: *neural_network*, *experience_buffer(MB)*, *gamma*

Ensure: *neural_network*

- 1: Sample mini-batch MB of experiences from buffer
- 2: **for** *s, a, s', r* ∈ MB **do** ▷ (state, action, next_state, reward)
- 3: target ← PREDICT(*neural_network*, *s*)
- 4: **if** *s'* is terminal **then**
- 5: target[*a*] = *r*
- 6: **else**
- 7: max-q = max(PREDICT(*neural_network*, *s'*)) ▷ max q-value from *s'*
- 8: target[*a*] = *r* + (max-q * *gamma*)
- 9: mse = MSE(*neural_net*, *state*, *target*)
- 10: gradient = COMPUTE_GRADIENT(*mse*, *neural_network*)
- 11: BACK_PROPAGATE_GRADIENT(*gradients*, *neural_network*)
- 12: **return** *neural_network*

Expected Results

Expected results

The plot on the right is the expected result. Notice that it has been obtained with an ϵ -greedy strategy, starting with $\epsilon = 1$ and multiplying it by 0.999 each epoch.

Seeding

Given the (particularly) high stochasticity of the method and the environment, for this lesson, we fixed a random seed equal to 15.




Figure: Note that obtaining this result requires time. You can stop the training after fewer iterations if you observe a growth in the reward.

Pedagogic Implementation

Today's lesson presents a simplified version of the code. In the next lessons, we will see a more efficient implementation that exploits Numpy, matrix multiplications, and advanced TensorFlow tools.

Number of Updates

By default, the suggested implementation performs the DQN update only once for episodes. A more efficient implementation exploits more iterations. However, this would require some tricks to avoid overfitting on the data (e.g., memory buffer shuffle).